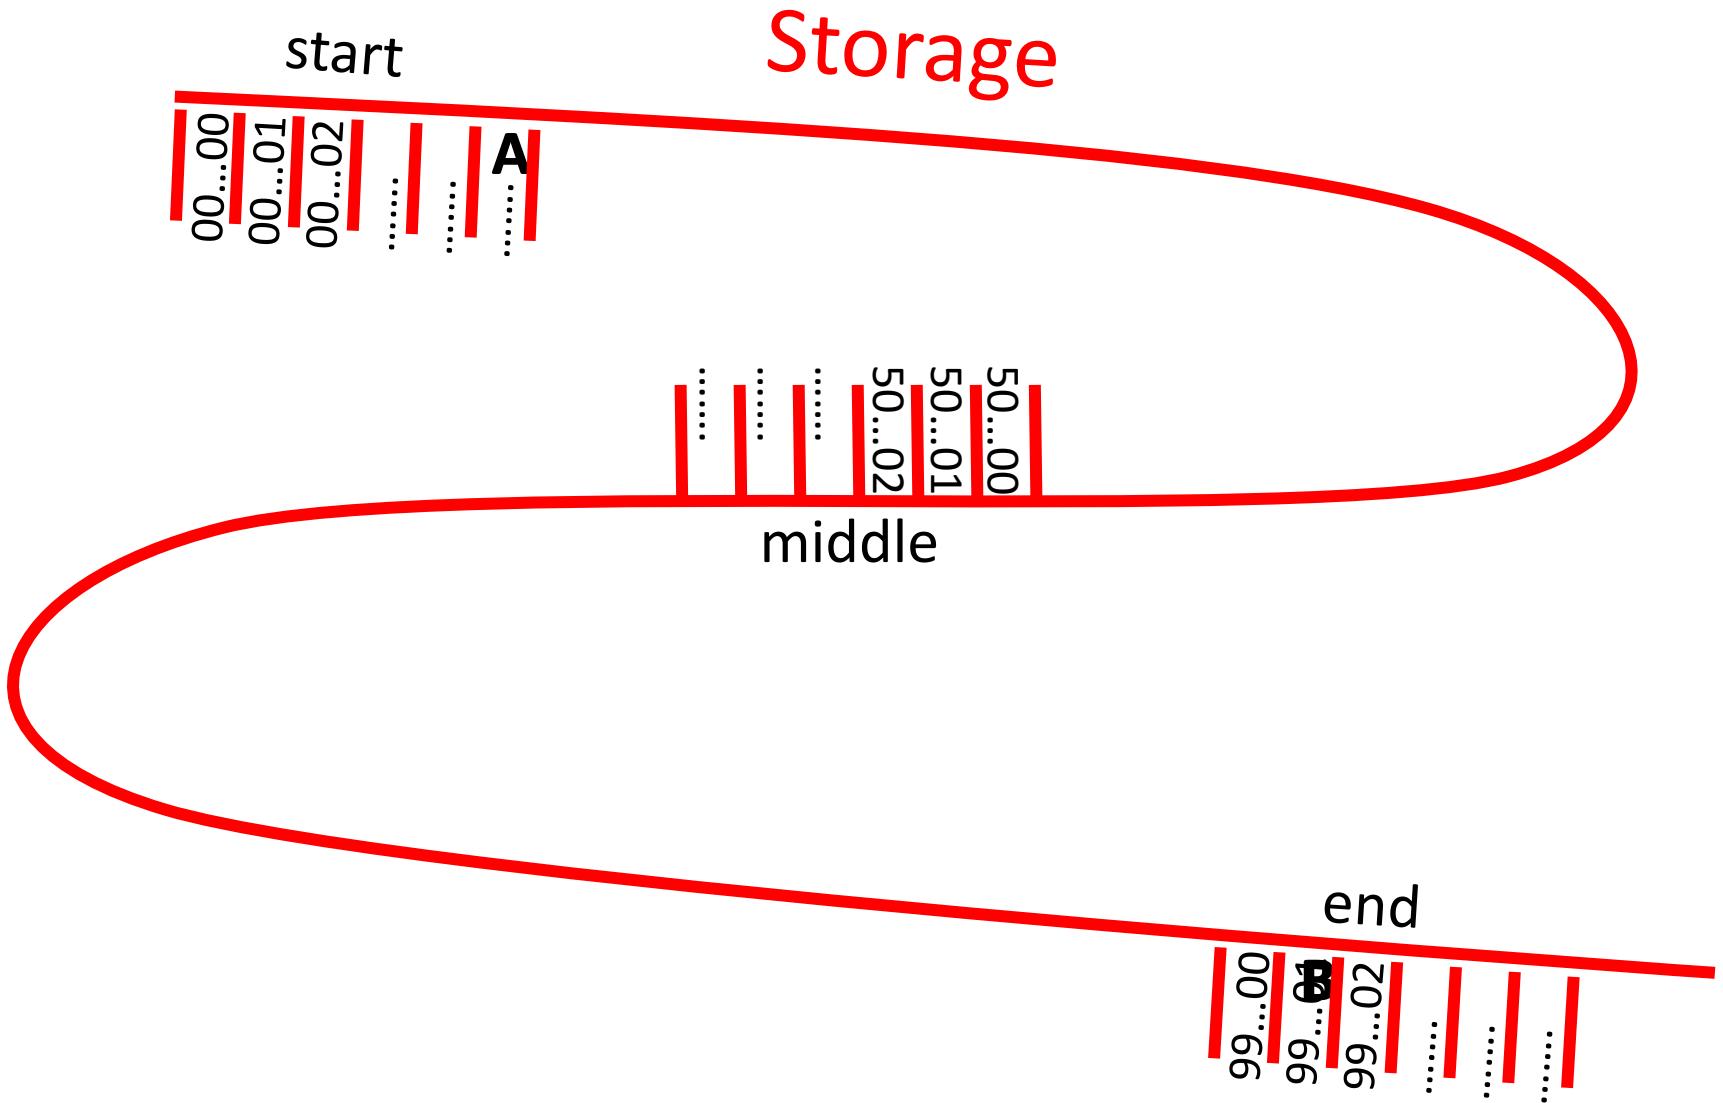
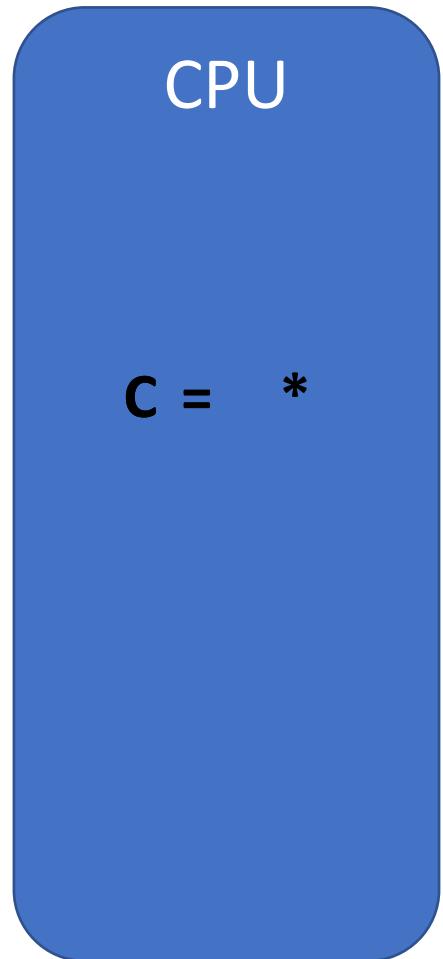


Storage Locality  
and counting words

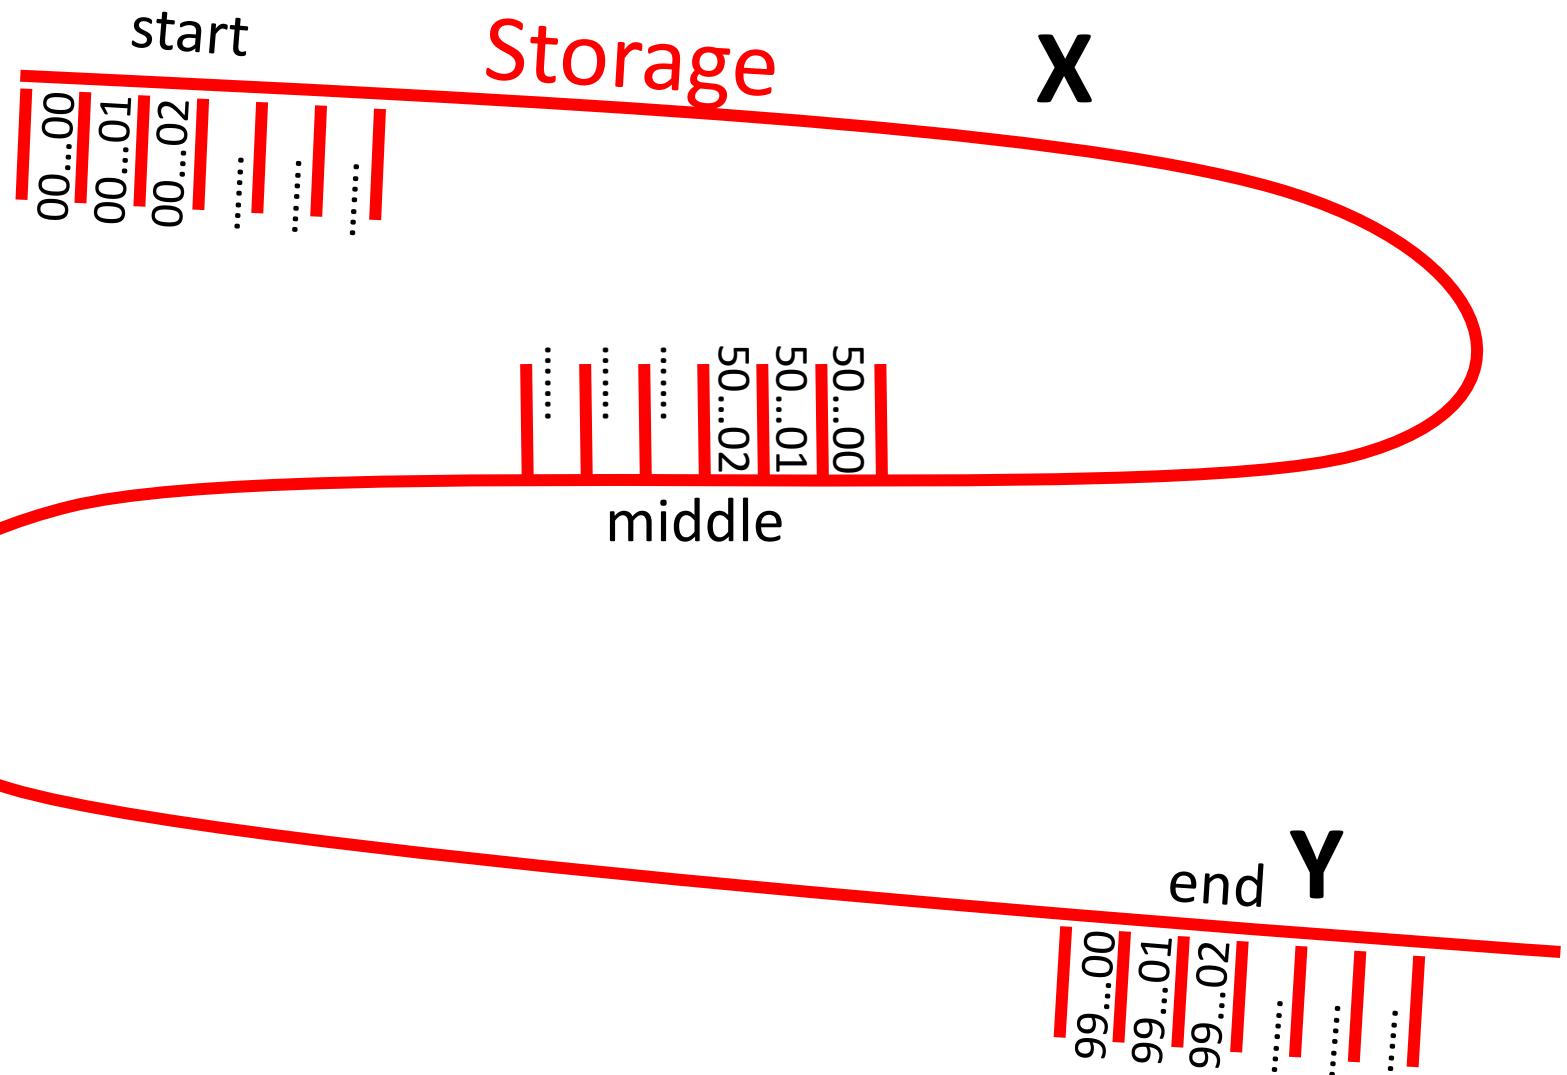
# NON-LOCAL STORAGE ACCESS



## LOCAL STORAGE ACCESS

CPU

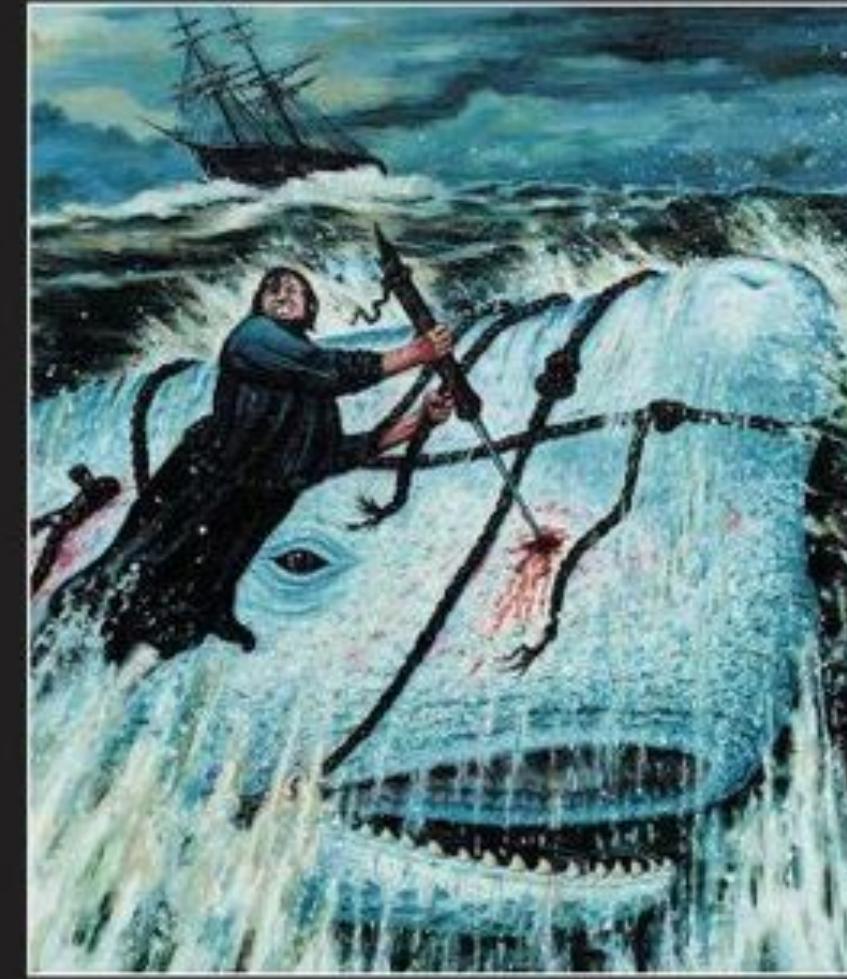
```
A=0  
For i in range(100000):  
    A+= X[i]  
  
For i in range(100000):  
    A-= Y[i]
```



# Counting the words in a long text

- 218718 words long
- 17150 different words
- How many times does each word occur?

Moby Dick  
HERMAN MELVILLE



**Task:** count the number of occurrences of each word in very long text.

- **Input:** Call me Ishmael. Some years ago--never mind how long precisely—having little or no money in my purse, and nothing particular to interest me onshore, I thought I would sail ...
- Moby dick: about 1.3MByte
- Desired output:
  - Call: 354
  - Me: 53423
  - Ismael: 1322
  - ....

# Simple solution

- Iterate over words. Update counter for current word.

```
In [3]: %%time
def simple_count(list):
    D={}
    for w in list:
        if w in D.keys():
            D[w]+=1
        else:
            D[w]=1
    return D
D=simple_count(all)
```

```
CPU times: user 56.9 s, sys: 335 ms, total: 57.2 s
Wall time: 57.1 s
```

# Lets use a sorted list

```
==== unsorted list:
```

```
the,vernacular,but,as,for,you,ye,carrion,rogues,turning,to,  
the,three,men,in,the,rigging,for,you,i,mean,to,mince,ye,up,  
for
```

```
==== sorted list:
```

5

```
lines,lingered,lingered,lingered,lingered,lingered,lingerin  
g,lingering,lingering,lingering,lingering,lingering,lingeri  
ng,lingering,lingers,lingo,lingo,lining,link,link,linked,li  
nk,linked,linked,links,links
```

8

# Sort-based solution

1. Sort words
2. Iterate over words. Update counter for current word.

In [5]:

```
%%time
def sort_count(list):
    S=sorted(list)
    D={}
    current=''
    count=0
    for w in S:
        if current==w:
            count+=1
        else:
            if current!='':
                D[current]=count
            count=1
            current=w
    return D
D=sort_count(all)
```

CPU times: user 180 ms, sys: 31.2 ms, total: 211 ms

Wall time: 190 ms

# Summary

- Sorting improves memory locality for word counting
- Improved memory locality reduces run-time
- Why? Because computer memory is organized in a hierarchy.

# Storage Latency

CPU

**C = \***

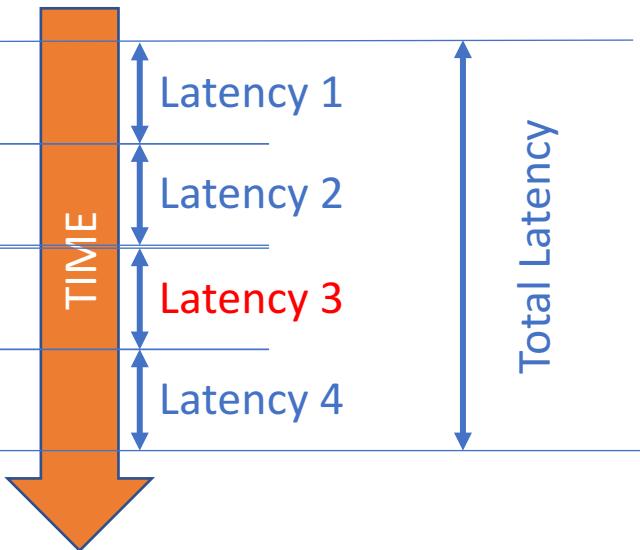


Storage



# Latencies

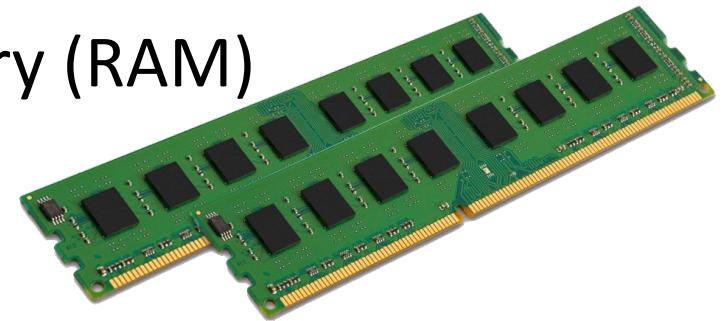
1. Read A
2. Read B
3.  $C=A*B$
4. Write C



With big data, most of the latency is memory latency (1,2,4), not computation (3)

# Storage Types

- Main Memory (RAM)



- Spinning disk



- Remote computer



# Summary

- The major source of latency in data analysis is reading and writing to storage
- Different types of storage offer different latency, capacity and price.
- Big data analytics revolves around methods for organizing storage and computation in ways that maximize speed while minimizing cost.
- Next, storage locality.

# Caches and the Memory Hierarchy

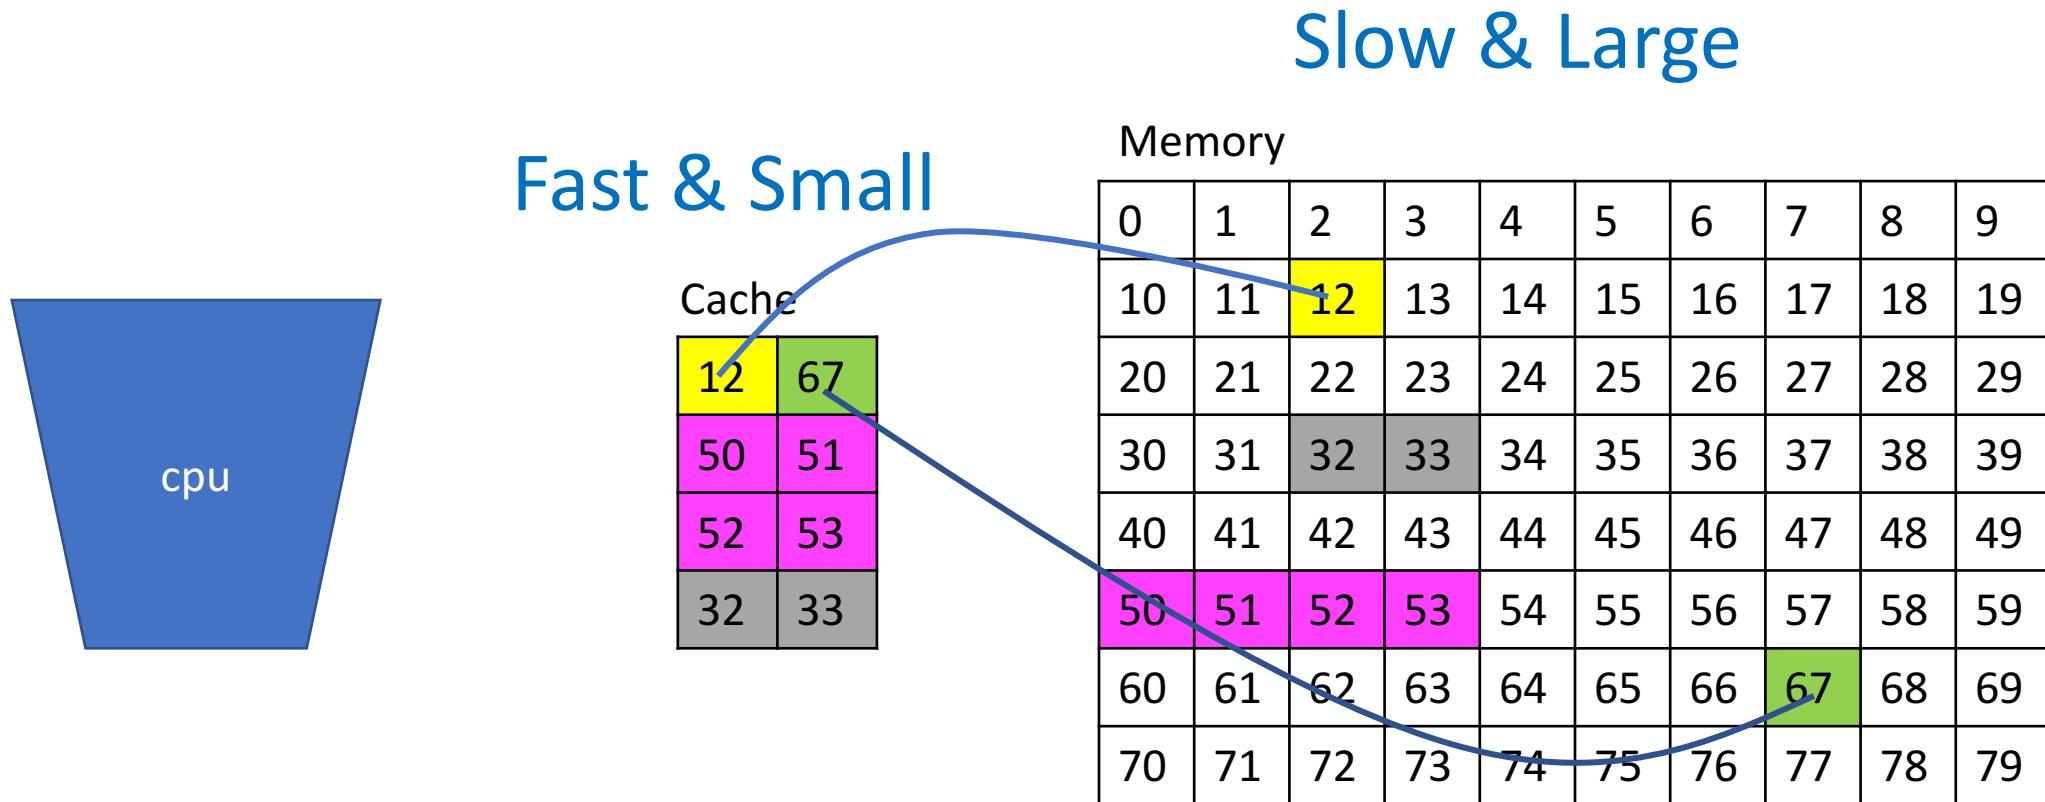
# Latency, size and price of computer memory

Given a budget, we need to trade off

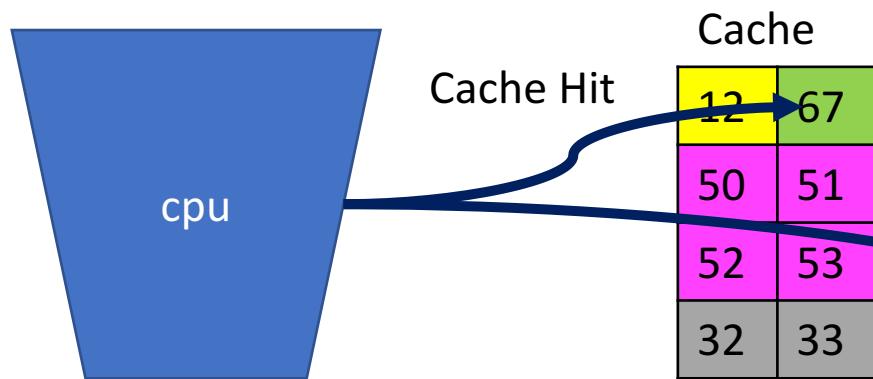
# \$10: Fast & Small


# \$10: Slow & Large

# Cache: The basic idea



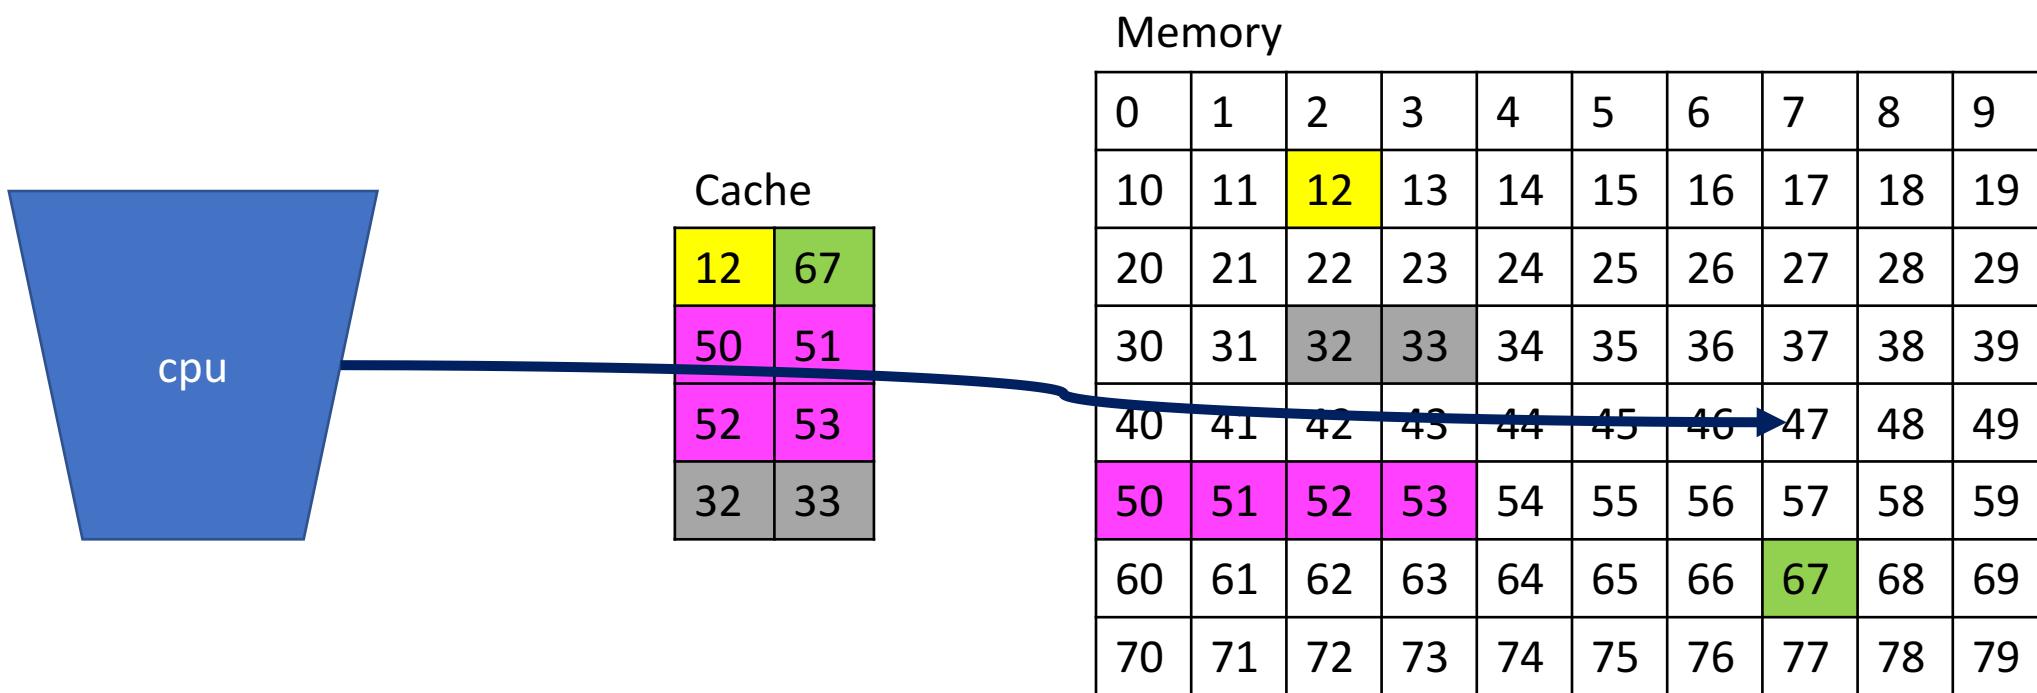
# Cache Hit



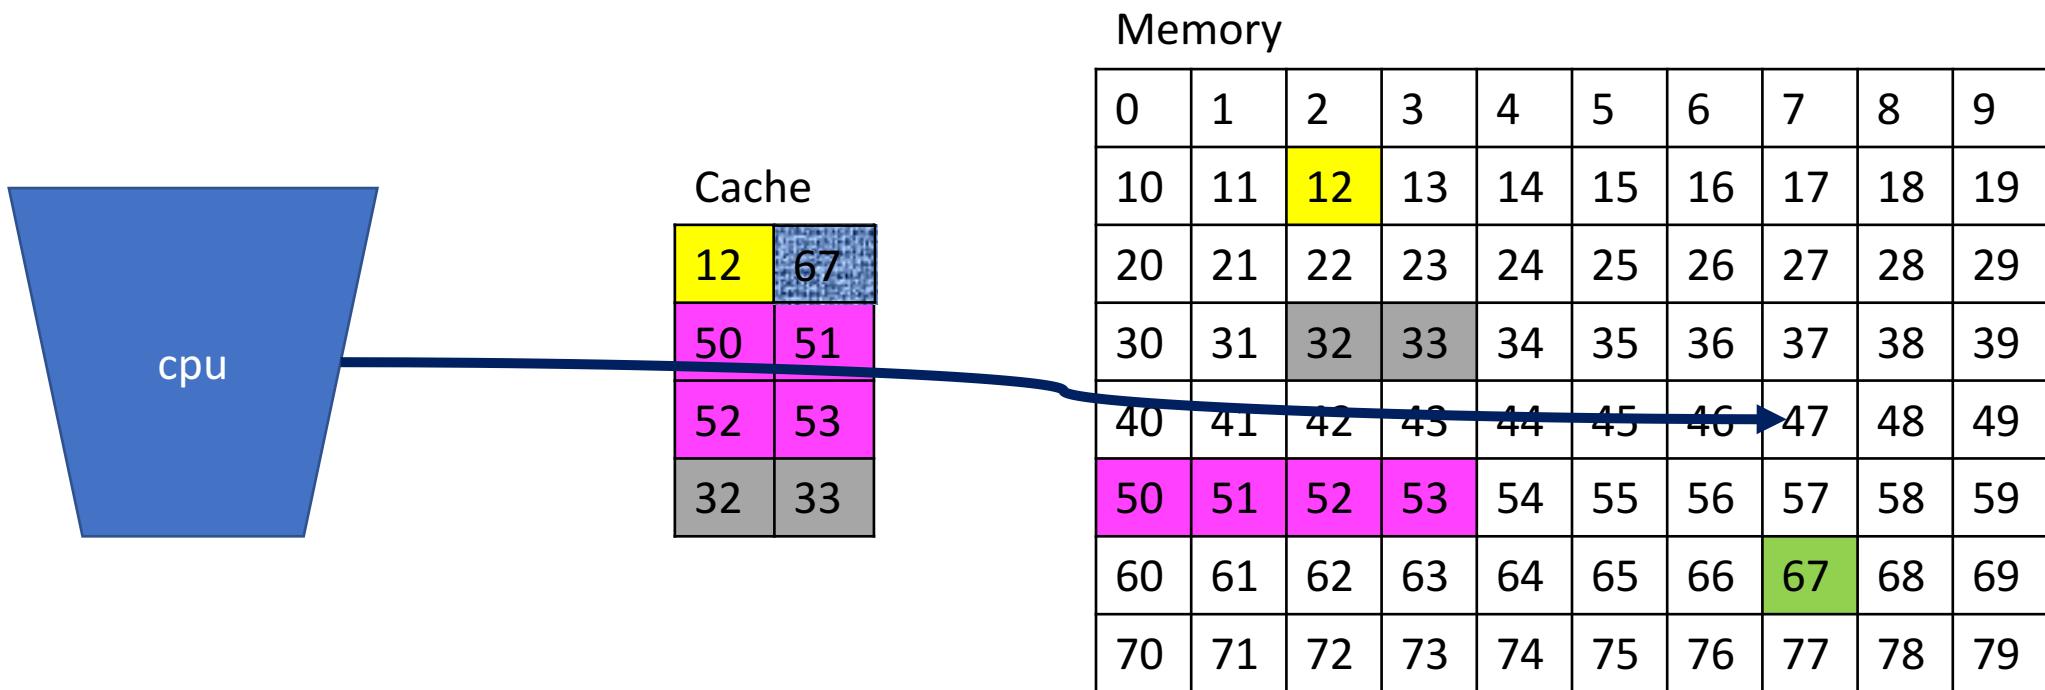
Memory

0	1	2	3	4	5	6	7	8	9
10	11	12	13	14	15	16	17	18	19
20	21	22	23	24	25	26	27	28	29
30	31	32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47	48	49
50	51	52	53	54	55	56	57	58	59
60	61	62	63	64	65	66	67	68	69
70	71	72	73	74	75	76	77	78	79

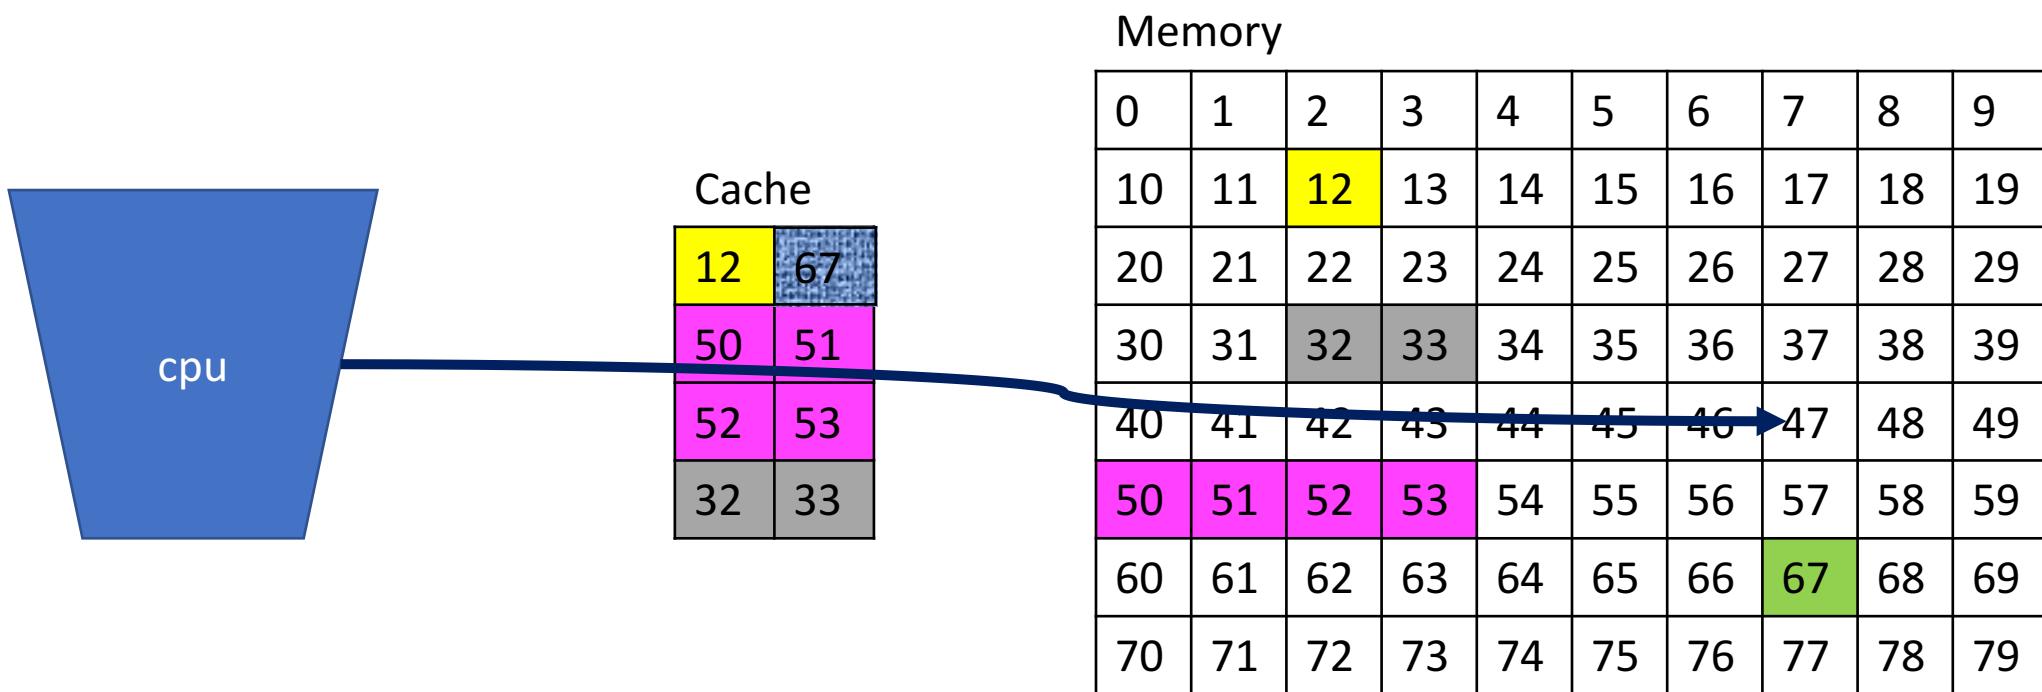
# Cache Miss



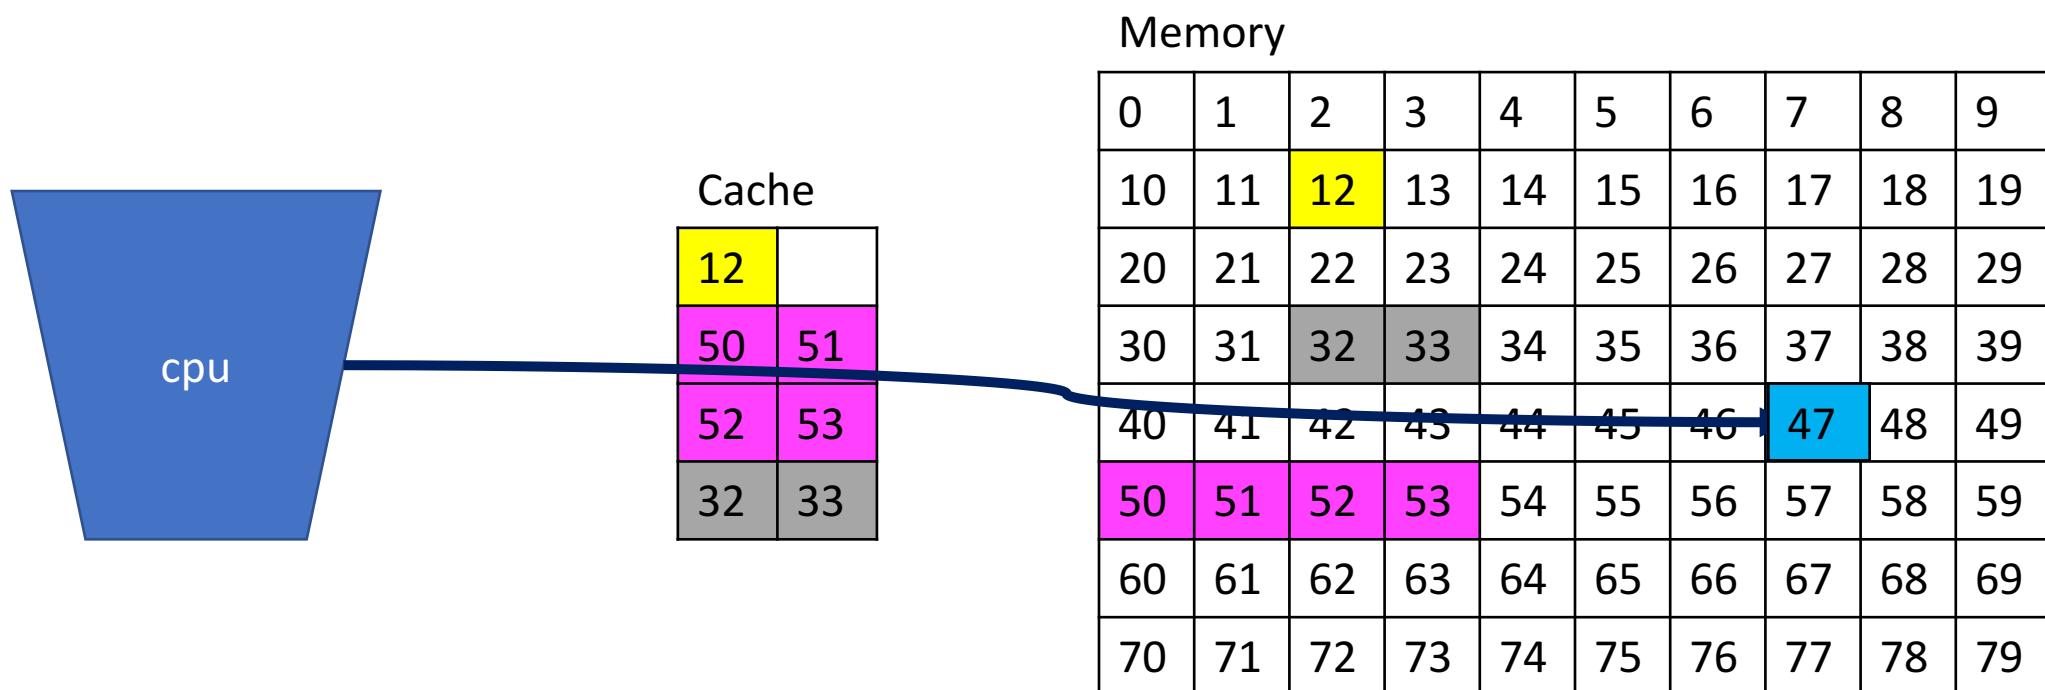
# Cache Miss Service: 1) Choose byte to drop



# Cache Miss Service: 2) write back



# Cache Miss Service: 3) Read In



# Access Locality

- The cache is effective If most accesses are hits.
  - Cache Hit Rate is high.
- **Temporal Locality:** Multiple accesses to **same** address within a short time period

# Spatial locality

- **Spatial Locality:** Multiple accesses to close-together addresses in short time period.
  - The difference between two sums.
  - Counting words by sorting
- Benefiting from spatial locality
  - Memory is partitioned into **Blocks/Lines** rather than single bytes.
  - Moving a block of memory takes much less time than moving each byte individually.
  - Memory locations that are close to each other are likely to fall in the same block.
  - Resulting in more cache hits.

# Unsorted word count / poor locality

```
==== unsorted list:
```

```
the, vernacular, but, as, for, you, ye, carrion, rogues, turning, to,
```

- Consider the memory access to the dictionary D:
- Count without sort:  
 $D[\text{the}]=12332, \dots, D[\text{but}]=943, \dots, D[\text{vernacular}]=10, \dots, D[\text{for}]=\dots$
- Temporal locality for very common words like “the”
- No spatial locality

# sorted word count / good locality

```
==== sorted list:
```

```
lines, lingered, lingered, lingered, lingered, lingered, lingered, lingerin  
g, lingering, lingering, lingering, lingering, lingering, lingeri  
ng, lingering, lingers, lingo, lingo, lining, link, link, linked, li  
ned, linked, linked, links, links
```

Entries to D are added one at a time.

1. D[lines]=33
2. D[lines]=33, D[lingered]=5
3. D[lines]=33, D[lingered]=5, D[lingering]=8

Assuming new entries are added at the end, this gives spatial locality.

Spatial locality makes code run much faster (X300)

# Summary

- Caching reduces storage latency by bringing relevant data close to the CPU.
- This requires that code exhibits access locality:
  - Temporal locality: Accessing the same location multiple times
  - Spatial locality: Accessing neighboring locations.

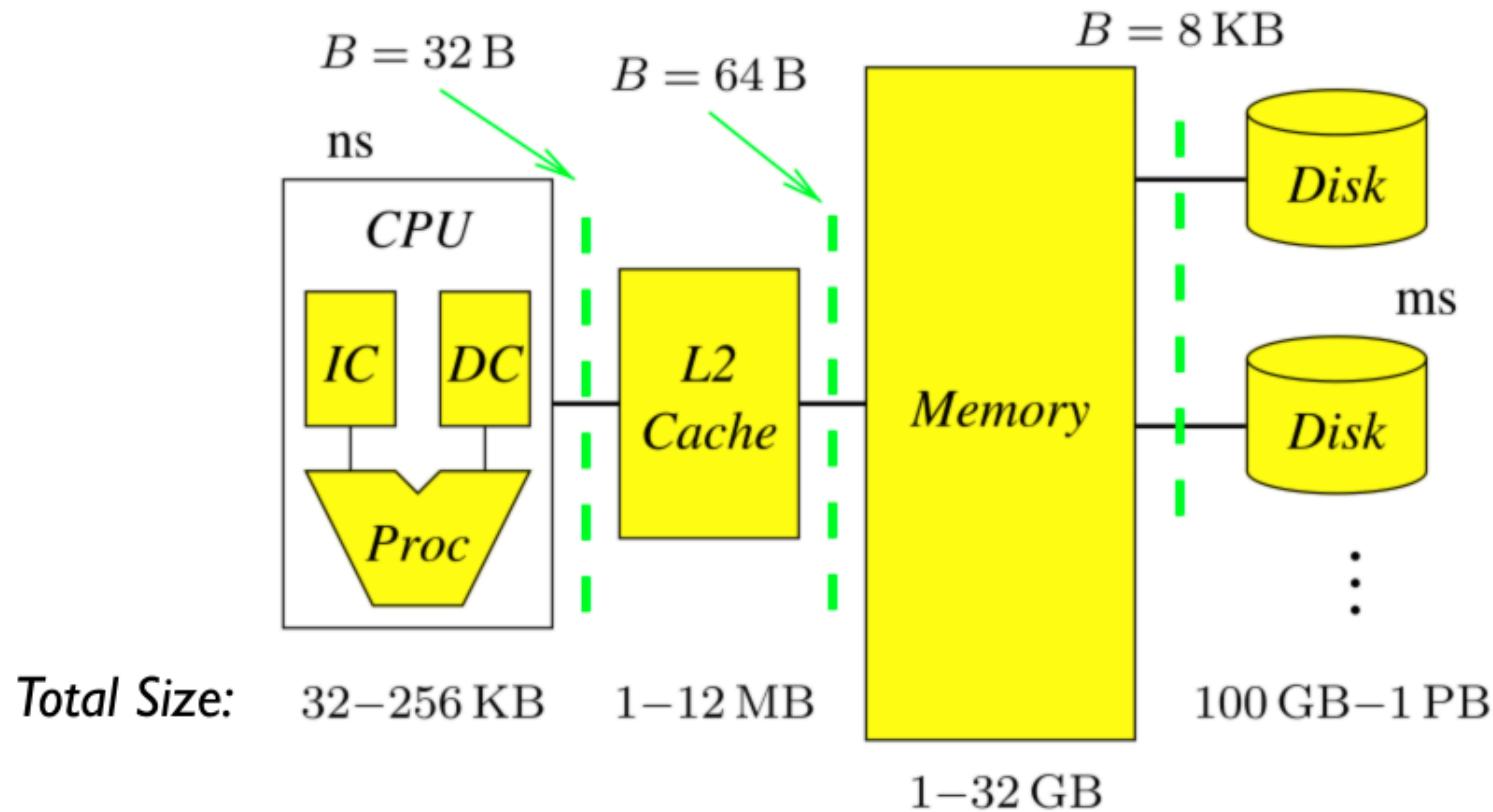
# The memory Hierarchy

# The Memory Hierarchy

- Real systems have several levels storage types:
  - Top of hierarchy: Small and fast storage close to CPU
  - Bottom of Hierarchy: Large and slow storage further from CPU
- Caching is used to transfer data between different levels of the hierarchy.
- To the programmer / compiler does not need to know
  - The hardware provides an **abstraction** : memory looks like a single large array.
- But performance depends on program's access pattern.

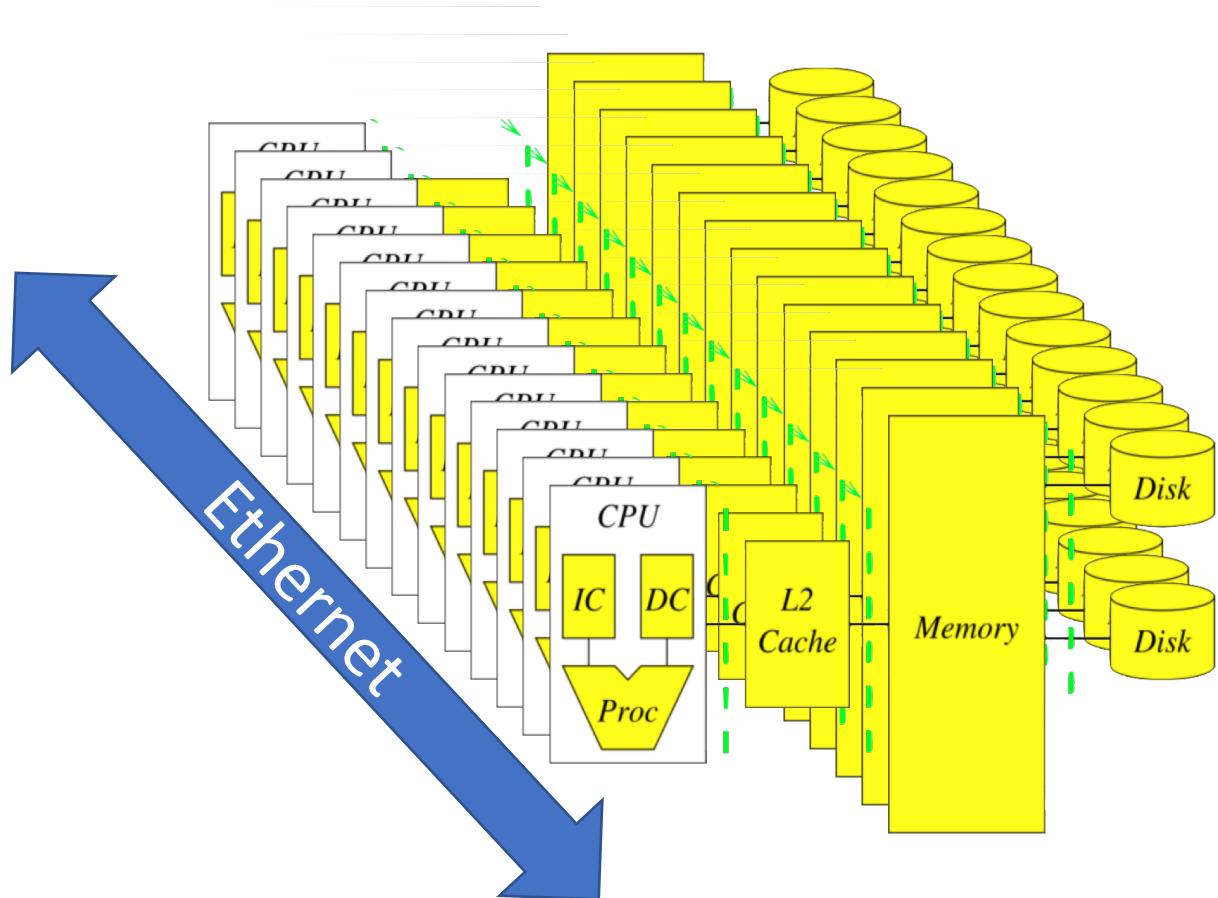
# The Memory Hierarchy

$B$ =Block size



# Computer clusters extend the memory hierarchy

- A data processing cluster is simply many computers linked through an ethernet connection.
- Storage is shared
- Locality: Data to reside on the computer will use it.
- “Caching” is replaced by “Shuffling”
- Abstraction is spark RDD.



# Sizes and latencies in a typical memory hierarchy.

	CPU (Registers)	L1 Cache	L2 Cache	L3 Cache	Main Memory	Disk Storage	Local Area Network
Size (bytes)	1KB	64KB	256KB	4MB	4-16GB	4-16TB	16TB - 1PB
Latency	300ps	1ns	5ns	20ns	100ns	2-10ms	2-10ms
Block size	64B	64B	64B	64B	32KB	64KB	1.5-64KB

Diagram illustrating the relationship between memory levels and their sizes and latencies:

- Size (bytes):** A red arrow points from CPU Registers to Main Memory, showing an increase in size across the hierarchy. The values are 1KB, 64KB, 256KB, 4MB, 4-16GB, 4-16TB, and 16TB - 1PB. A red box highlights the range from 16TB to 1PB, labeled "12 orders of magnitude".
- Latency:** A red arrow points from CPU Registers to Disk Storage, showing an increase in latency. The values are 300ps, 1ns, 5ns, 20ns, 100ns, 2-10ms, and 2-10ms. A red box highlights the range from 2-10ms, labeled "6 orders of magnitude".
- Block size:** A red arrow points from CPU Registers to Local Area Network, showing a general trend of increasing block size. The values are 64B, 64B, 64B, 64B, 32KB, 64KB, and 1.5-64KB. A red box highlights the range from 1.5-64KB.

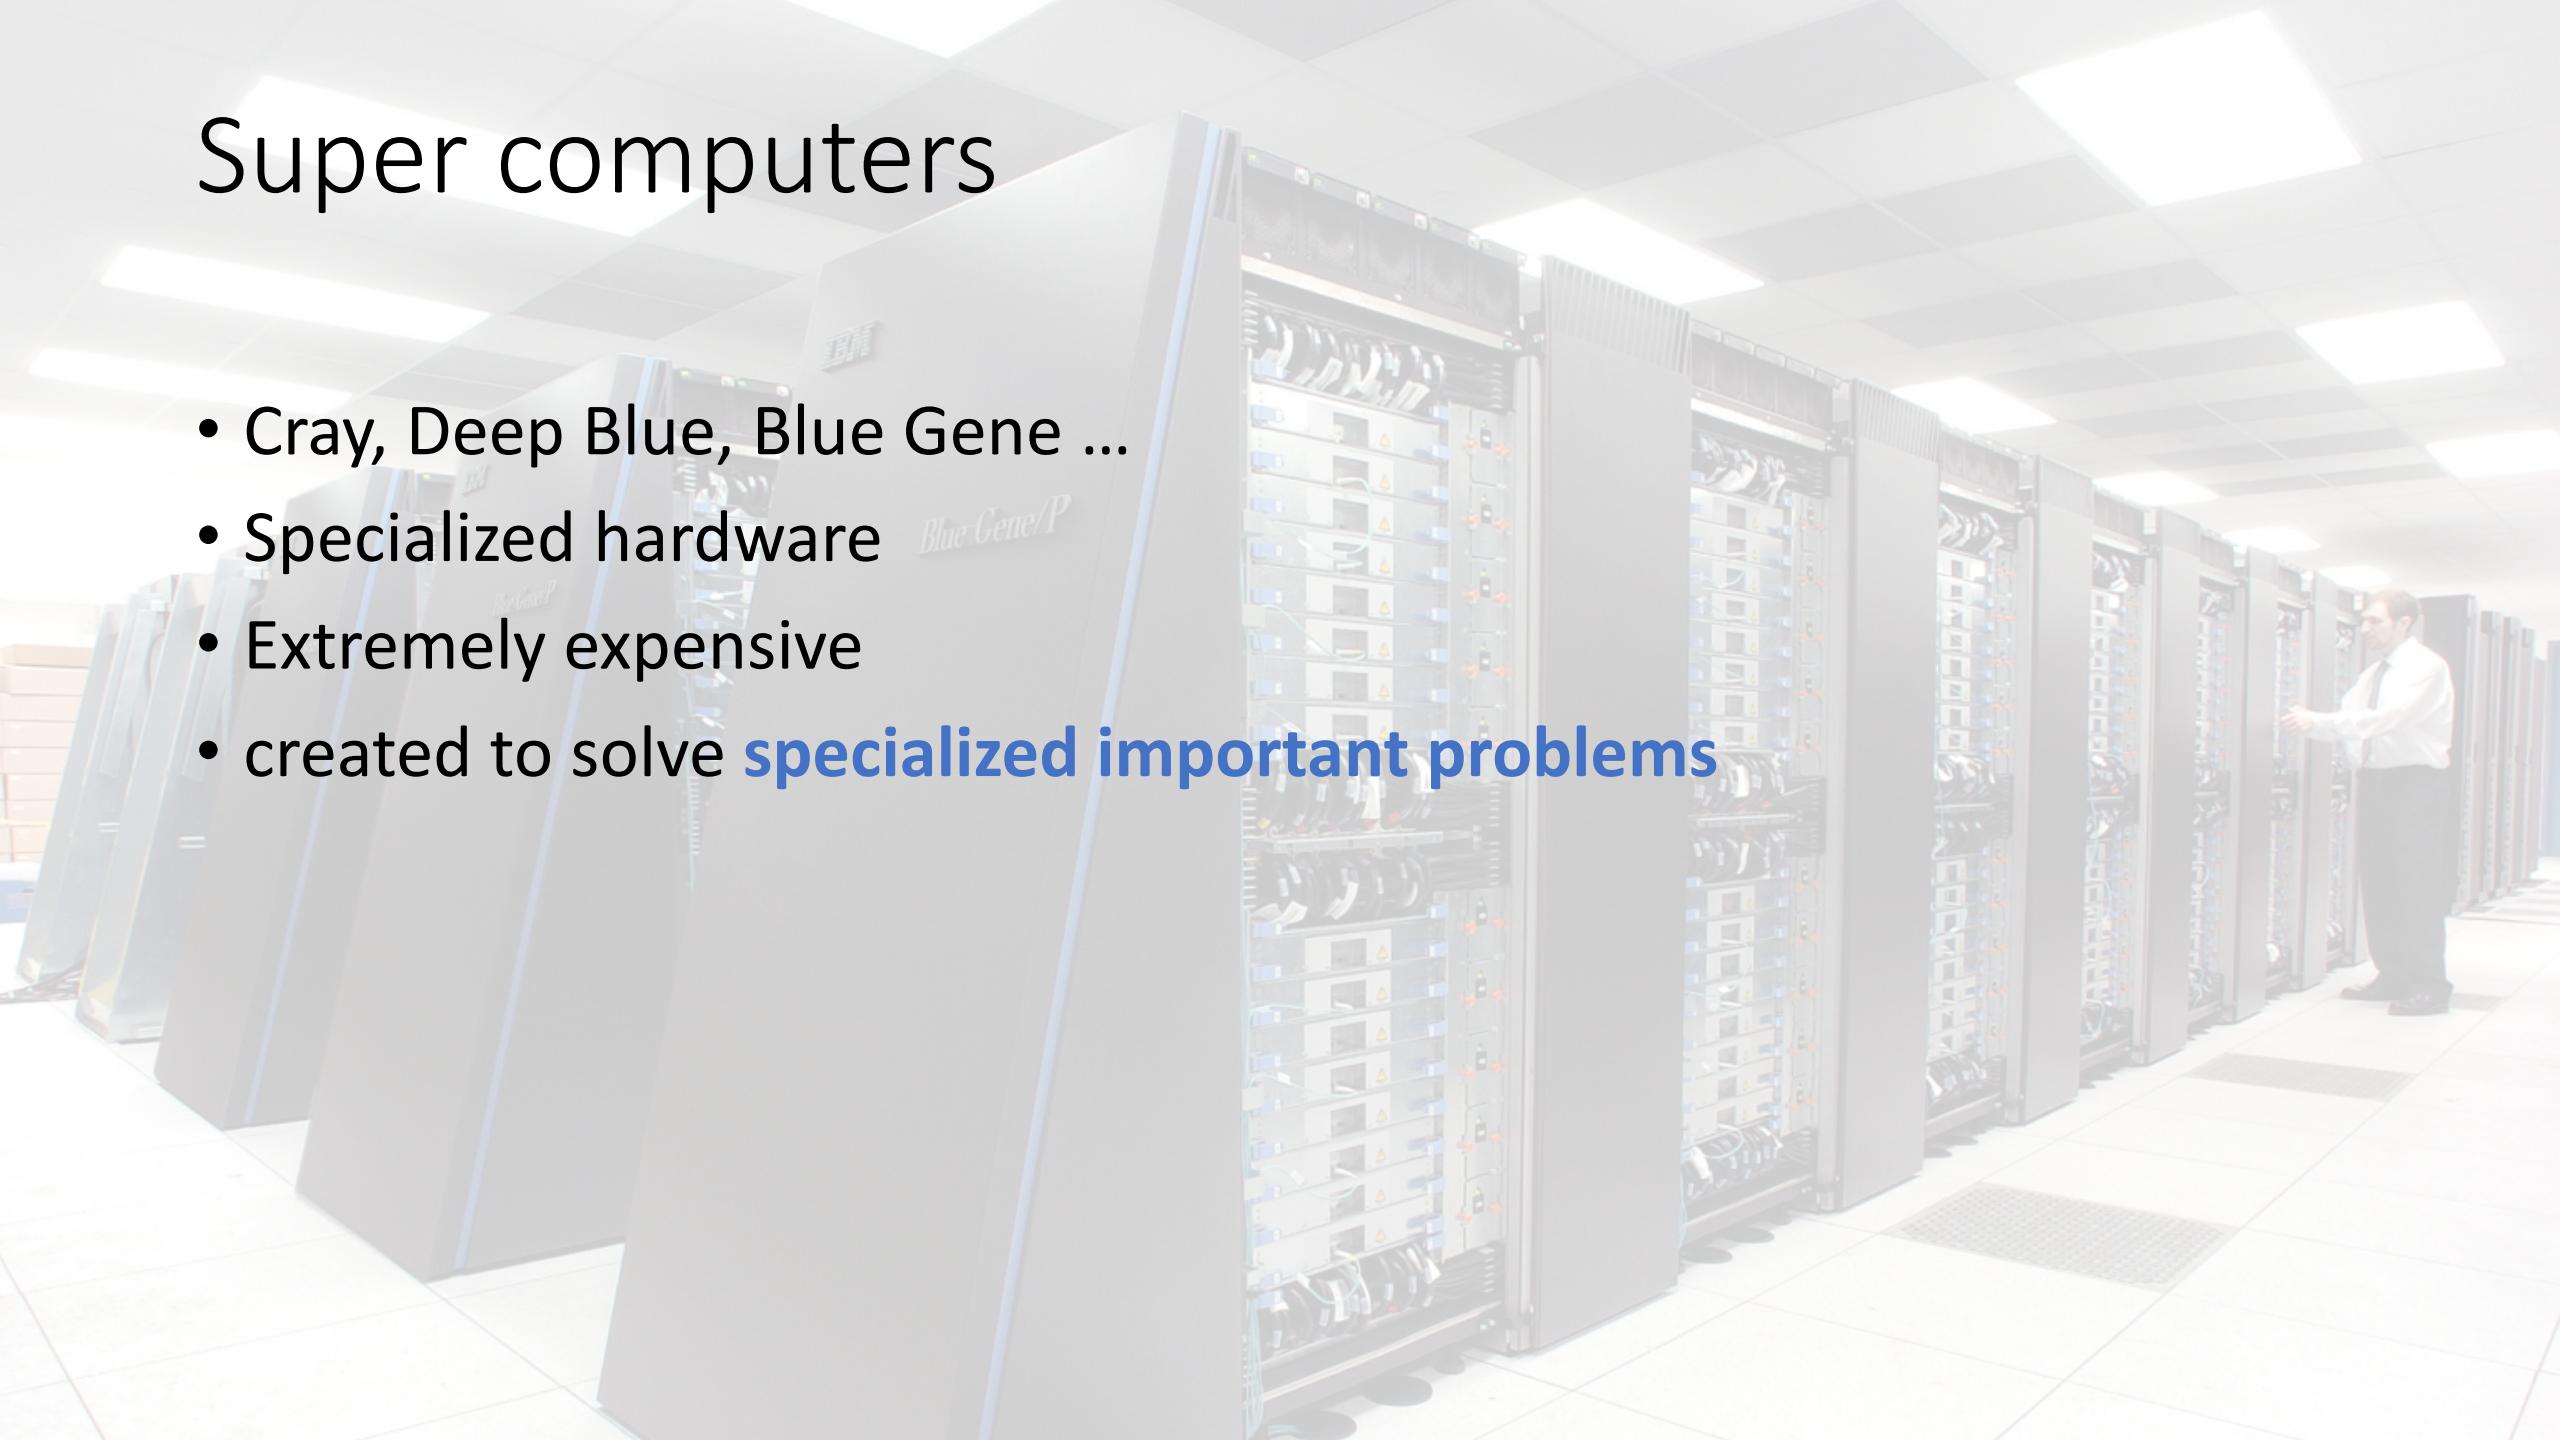
# Summary

- Memory Hierarchy: combining storage banks with different latencies.
- Clusters: multiple computers, connected by ethernet, that share their storage.

A short history of affordable  
massive computing.

# Super computers

- Cray, Deep Blue, Blue Gene ...
- Specialized hardware
- Extremely expensive
- created to solve **specialized important problems**



# Data Centers



# Data Centers

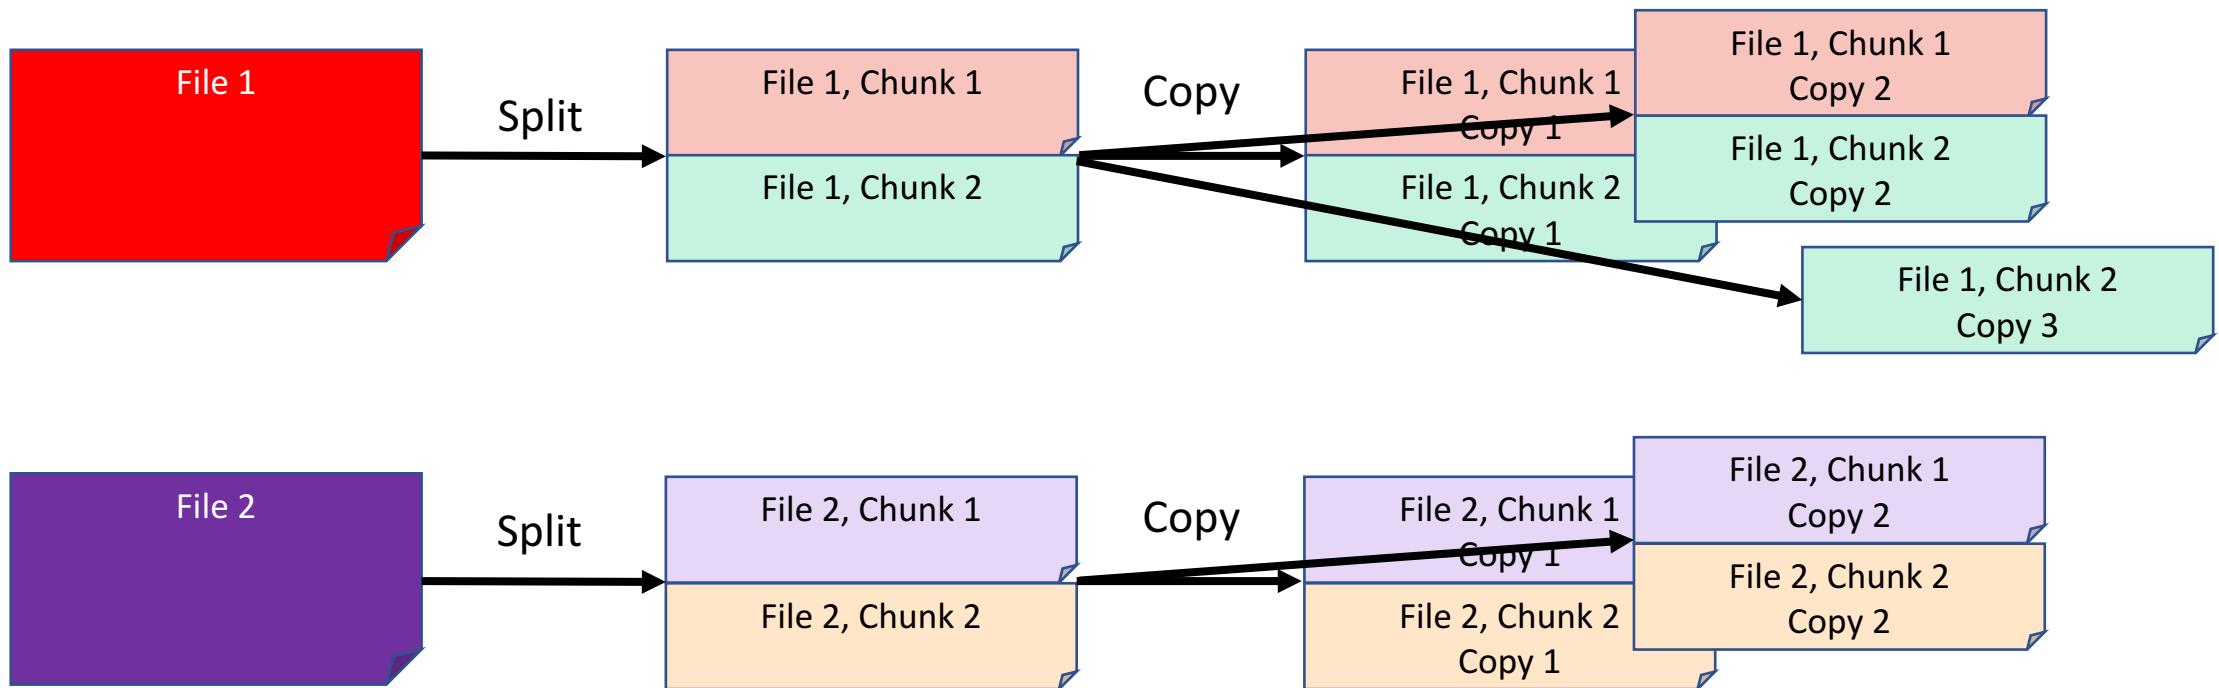
- The physical aspect of "the cloud"
- Collection of commodity computers
- VAST number of computers (100,000's)
- Created to provide computation for large and small organizations.
- Computation as a commodity.



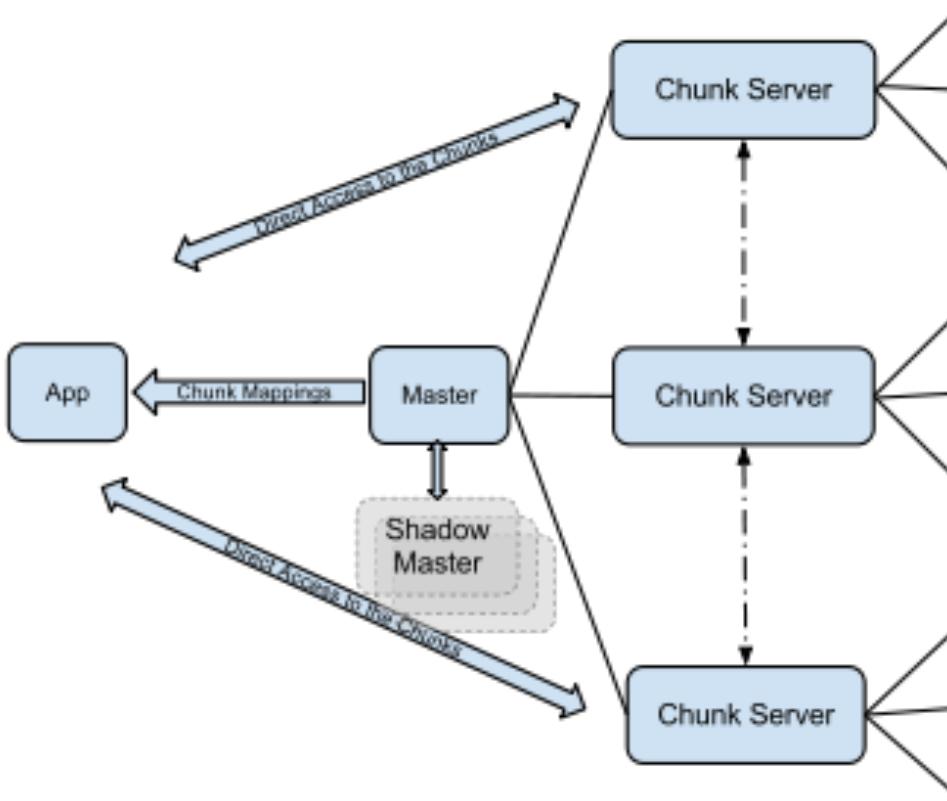
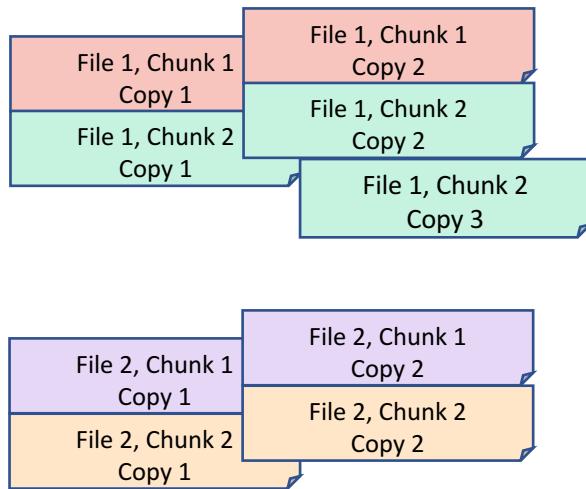
# Making History: Google 2003

- Larry Page and Sergey Brin develop a method for storing very large files on multiple **commodity** computers.
- Each file is broken into fixed-size **chunks**.
- Each chunk is stored on multiple **chunk servers**.
- The locations of the chunks is managed by the **master**

# HDFS: Chunking files



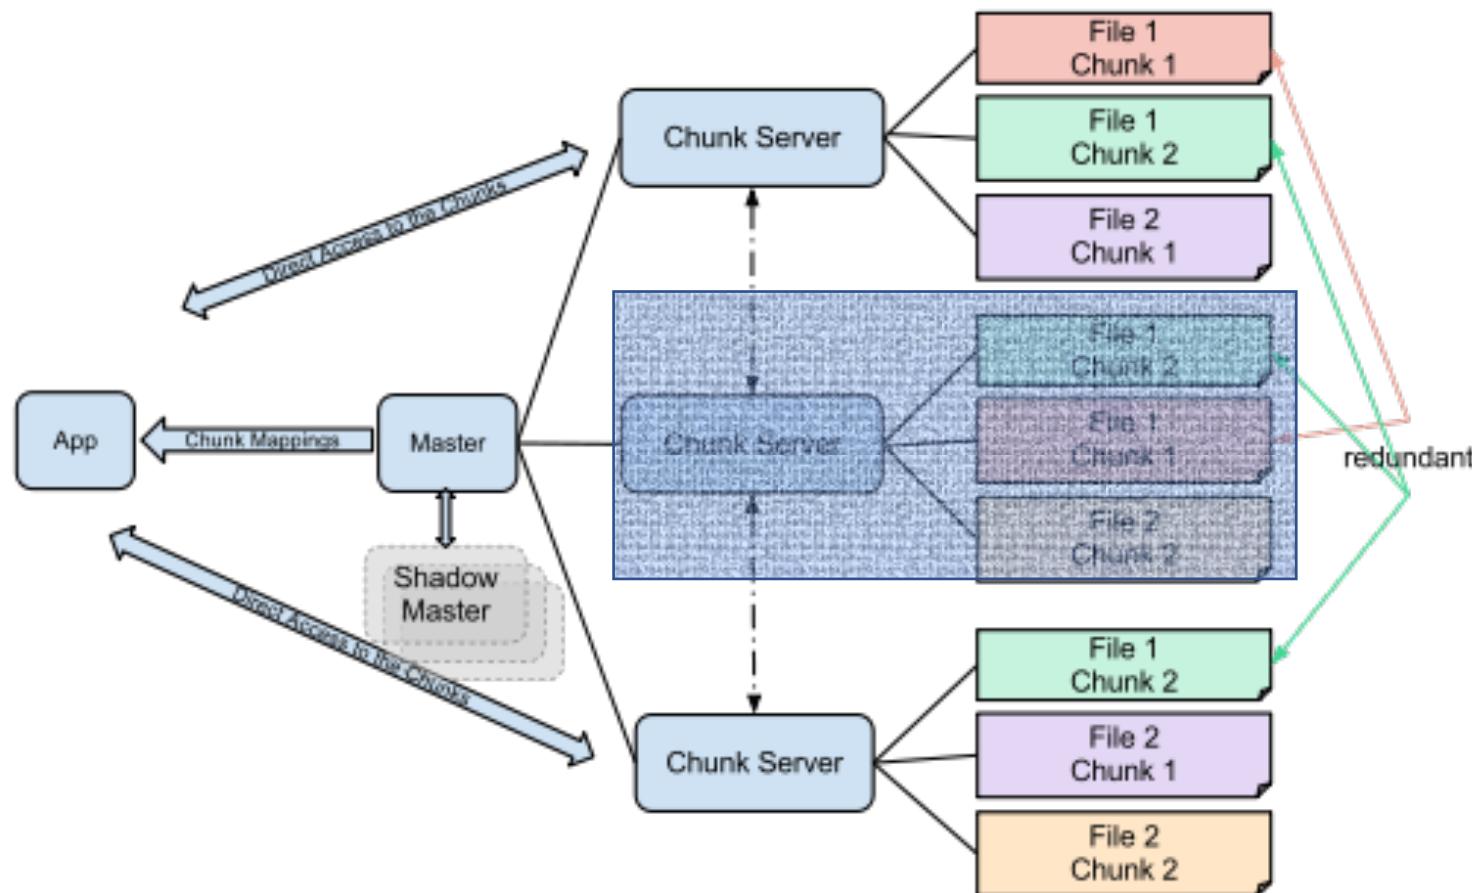
# HDFS: Distributing Chunks



# Properties of GFS/HDFS

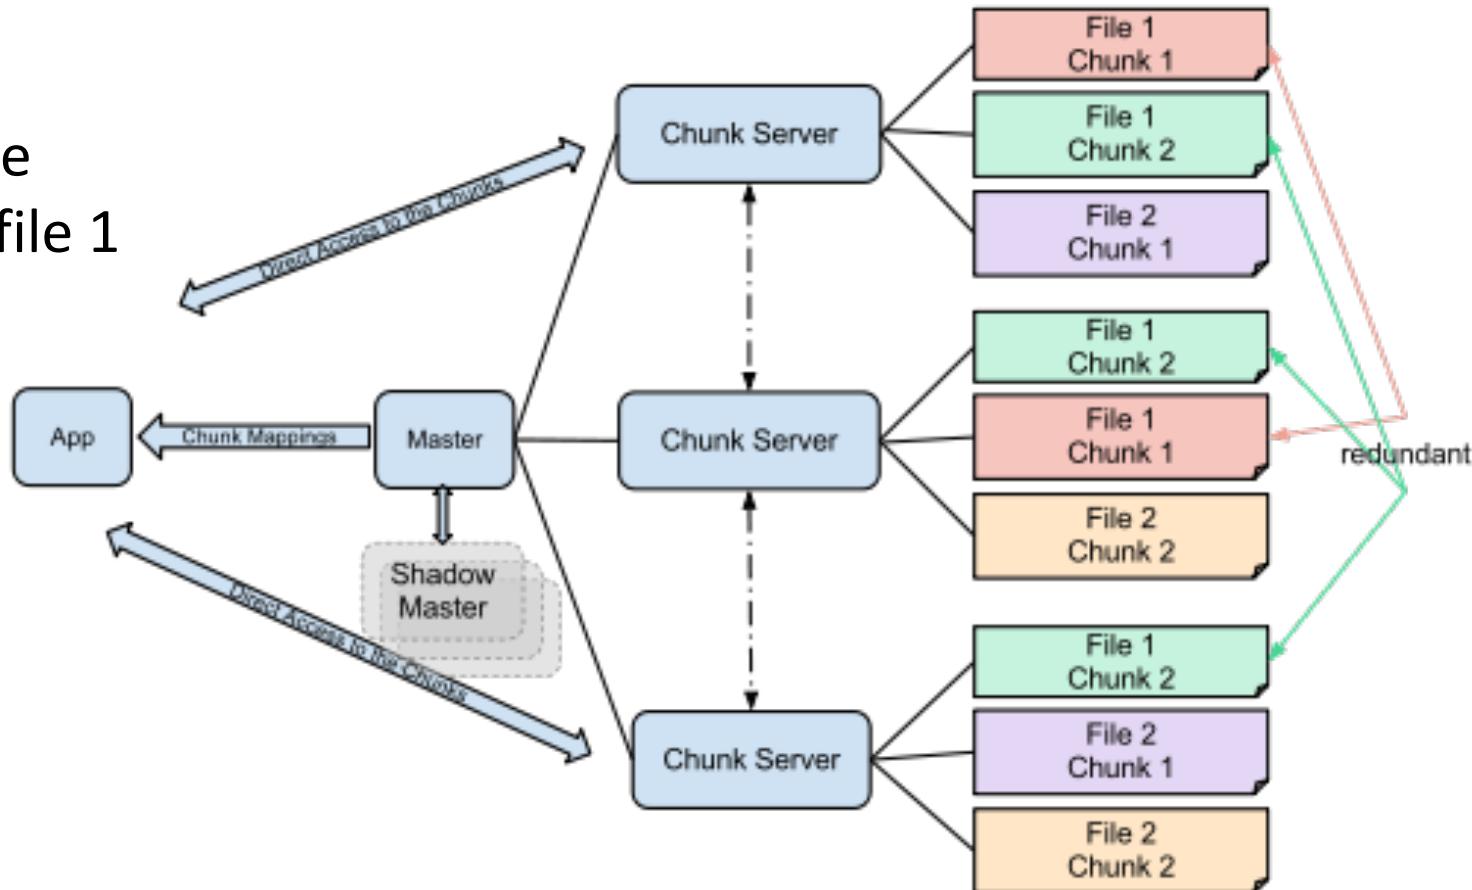
- **Commodity Hardware:** Low cost per byte of storage.
- **Locality:** data stored close to CPU.
- **Redundancy:** can recover from server failures.
- **Simple abstraction:** looks to user like standard file system (files, directories, etc.) Chunk mechanism is hidden.

# Redundancy



# Locality

Task:  
Sum all of the  
elements in file 1



# Map-Reduce

- HDFS is a **storage abstraction**
- **Map-Reduce** is a **computation abstraction** that works well with HDFS
- Allows programmer to specify parallel computation without knowing how the hardware is organized.
- We will describe Map-Reduce, using Spark, in a later section.

# Spark

- Developed by Matei Zaharia , amplab, 2014
- Hadoop uses shared **file system** (disk)
- Spark uses shared **memory** – faster, lower latency.
- Will be used in this course
  
- Recall word count by sorting,  
we will redo it using map-reduce!

# Summary

- Big data analysis is performed on large clusters of commodity computers.
- HDFS (Hadoop file system): break down files to chunks, make copies, distribute randomly.
- Hadoop Map-Reduce: a computation abstraction that works well with HDFS
- Spark: Sharing **memory** instead of sharing **disk**.